

Randomized Algorithms (Ch13)

- Randomization
- Contention Resolution
- Global Min-cut
- Linearity of Expectation
- Maximum 3-Satisfiability
- Derandomization
- Monte Carlo and Las Vegas Algorithms
- Universal Hashing
- Closest Pair Problem
- Cache Maintenance Problem
- Chernoff Bounds
- Load Balancing

Randomization

- Algorithmic design approaches

Greedy

Divide-and-conquer

Dynamic programming

Network flow

Randomization

- **Randomization:** Allow fair coin flip (access to a pseudo-random number generator in practice) in unit time to decide the next step.
- Randomization can lead to simplest, fastest, or only known algorithm for a particular problem. Examples: Symmetry-breaking protocols, graph algorithms, quicksort, hashing, load balancing,...

Preliminaries for randomization

Experiments and outcomes

Experiment	Outcome
Toss a coin	$\{\text{heads}, \text{tails}\} = \{\text{H}, \text{T}\}$
Roll a dice	$\{1, 2, 3, 4, 5, 6\}$
Roll two dices	$\{1, \dots, 6\} \times \{1, \dots, 6\}$ or $A \subseteq \{1, \dots, 6\}, A \leq 2$
Buy 3 lottery tickets (out of 100,000)	3-element subsets of $\{1, \dots, 100,000\}$



Sample space and events

- Set S of all outcomes of an experiment is called the **sample space**.
- Any subset A of S is called an **event**.
- Examples

Experiment	Sample space	event
Toss 2 coins	$S = \{\text{H,T}\} \times \{\text{H,T}\}$	get exactly 1 H $A = \{(\text{H,T}), (\text{T,H})\}$
Roll 2 dice	$S = \{1, \dots, 6\} \times \{1, \dots, 6\}$	sum of dice is 6 $A = \{(1, 5), (2, 4), (3, 3), (4, 2), (5, 1)\}$

Probability

- **Equal likelihood assumption:** In an experiment, each of all possible outcomes has the **same likelihood** of occurrence, or the same probability of occurrence.
- Under the equal likelihood assumption, let S be the sample space for an experiment. For $a \in S$ and $A \subseteq S$,

$$\Pr[\{a\}] = \Pr[a] = \frac{1}{|S|} \quad \text{probability that } a \text{ occurs}$$

$$\Pr[A] = \frac{|A|}{|S|} \quad \text{probability } A \text{ occurs}$$

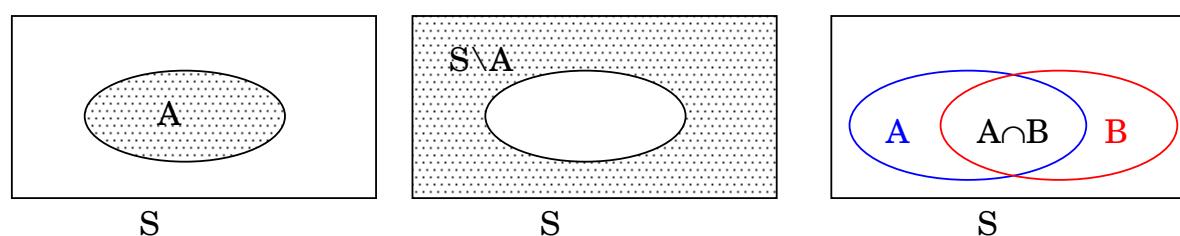
- **Examples**
 - **Probability of getting H (heads) in tossing a coin:**
Sample space $S = \{\text{H,T}\}$, event $A = \{\text{H}\}$, $\Pr[A] = \frac{|A|}{|S|} = \frac{1}{2}$.
 - **Probability of getting even number in rolling a dice:**

$$S = \{1, 2, 3, 4, 5, 6\}, \text{event } A = \{2, 4, 6\}, \Pr[A] = \frac{|A|}{|S|} = \frac{3}{6} = \frac{1}{2}.$$

- Probability is a measure quantifying the **likelihood** that an event A in a sample space S will occur, a function $\Pr : S \rightarrow [0, 1]$ s.t. for any $A \subseteq S$, $\Pr[A] = \sum_{a \in A} \Pr[a]$ satisfying:
 $\Pr[\emptyset] = 0$, $\Pr[S] = 1$, $0 \leq \Pr[A] \leq 1$ for all $A \subseteq S$ and
 $\Pr[A \cup B] = \Pr[A] + \Pr[B]$ for any disjoint $A, B \subseteq S$.
- Example: $S = \{a_1, a_2, \dots, a_n\}$, $\Pr[a_i] = p_i \geq 0$, $\sum_{i=1}^n p_i = 1$.
For $A \subseteq S$, $\Pr[A] = \sum_{a_i \in A} p_i$.
Equal likelihood assumption: $\Pr[a_i] = 1/n$ for every $1 \leq i \leq n$.
- For a sample space S and events $A, B \subseteq S$,
 $\Pr[\overline{A}] = 1 - \Pr[A]$, where $\overline{A} = S \setminus A$.

Union Bound

$$\Pr[A \cup B] = \Pr[A] + \Pr[B] - \Pr[A \cap B] \leq \Pr[A] + \Pr[B].$$



Discrete random variable

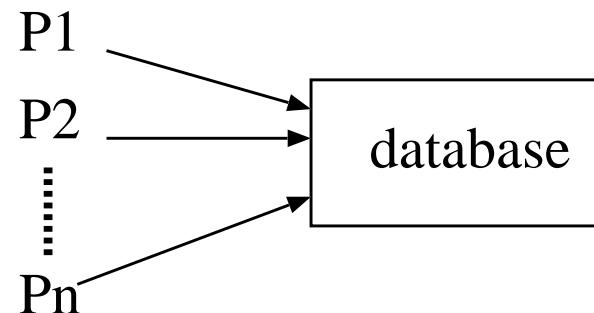
- A variable that takes discrete values with certain probability, (e.g., $X = 0, 1, \dots, n$ with $\Pr[X = i] = p_i$).
- Example, X = amount of money one wins by buying 1 lottery ticket:
There are 1000 tickets, 1 wins \$10,000, 10 wins \$100, the rest wins nothing;
 $\Pr[X = 10,000] = \frac{1}{1000}$, $\Pr[X = 100] = \frac{10}{1000}$, $\Pr[X = 0] = \frac{989}{1000}$.
- Let X be a discrete random variable taking values v_1, \dots, v_k . The expectation of X is $E[X] = v_1 \Pr[X = v_1] + \dots + v_k \Pr[X = v_k]$.

Example:

$$\begin{aligned}
 E[X] &= E[\text{money one wins}] \\
 &= 10000 \Pr[X = 10000] + 100 \Pr[X = 100] + 0 \Pr[X = 0] \\
 &= 10000 \frac{1}{1000} + 100 \frac{10}{1000} = 11.
 \end{aligned}$$

Contention Resolution

- **Contention resolution in a distributed system.**
 - **Instance:** Given n processes P_1, \dots, P_n , each competing for access to a shared database. If two or more processes access the database at the same time, all processes are locked out.
 - **Object:** A protocol to ensure all processes get through on a regular basis.
 - **Restriction:** Processes can not communicate.
- **Challenge:** Need a symmetry-breaking paradigm.



- A randomized protocol for contention resolution

Each process requests access to the database at time t with probability

$$p = 1/n.$$

- **Claim 1:** Let $S(i, t)$ be the event that process i succeeds in accessing the database at time t . Then $\frac{1}{en} \leq \Pr[S(i, t)] \leq \frac{1}{2n}$.

Proof. By independence, $\Pr[S(i, t)] = p(1 - p)^{n-1} = \frac{1}{n}(1 - \frac{1}{n})^{n-1}$. Since $\frac{1}{e} \leq (1 - \frac{1}{n})^{n-1} \leq \frac{1}{2}$, the claim holds. \square

- **Useful facts from calculus.** As n increases from 2,

$$\lim_{n \rightarrow \infty} (1 + \frac{1}{n})^n = e, e = 2.718\dots$$

$(1 - \frac{1}{n})^n$ converges monotonically from $1/4$ up to $1/e$.

$(1 - \frac{1}{n})^{n-1}$ converges monotonically from $1/2$ down to $1/e$.

- **Claim 2: The probability that process i fails to access the database in en rounds is at most $1/e$; after $en(c \ln n)$ rounds, the probability is at most n^{-c} .**

Proof. Let $F(i, t)$ be the event that process i fails to access the database in rounds 1 through t . By independence and Claim 1, $\Pr[F(i, t)] \leq (1 - \frac{1}{en})^t$.

Choose $t = \lceil en \rceil$,

$$\Pr[F(i, t)] \leq (1 - \frac{1}{en})^{\lceil en \rceil} \leq (1 - \frac{1}{en})^{en} \leq \frac{1}{e}.$$

Choose $t = \lceil en \rceil \lceil c \ln n \rceil$,

$$\Pr[F(i, t)] \leq (\frac{1}{e})^{c \ln n} = n^{-c}.$$

□

- **Claim 3: The probability that all processes succeed within $2en \ln n$ rounds is at least $1 - 1/n$.**

Proof. Let $F(t)$ be the event that at least one of the n processes fails to access the database in any of the rounds 1 through t . By **Union Bound** (Given events E_1, \dots, E_n , $\Pr[\cup_{i=1}^n E_i] \leq \sum_{i=1}^n \Pr[E_i]$; taking $E_i = F(i, t)$) and Claim 2,

$$\Pr[F(t)] = \Pr[\cup_{i=1}^n F(i, t)] \leq \sum_{i=1}^n \Pr[F(i, t)] \leq n\left(1 - \frac{1}{en}\right)^t.$$

Choose $t = 2 \lceil en \rceil \lceil \ln n \rceil$ gives $\Pr[F(t)] \leq nn^{-2} = 1/n$. □

Global Min-Cut

- **Global min-cut problem:**

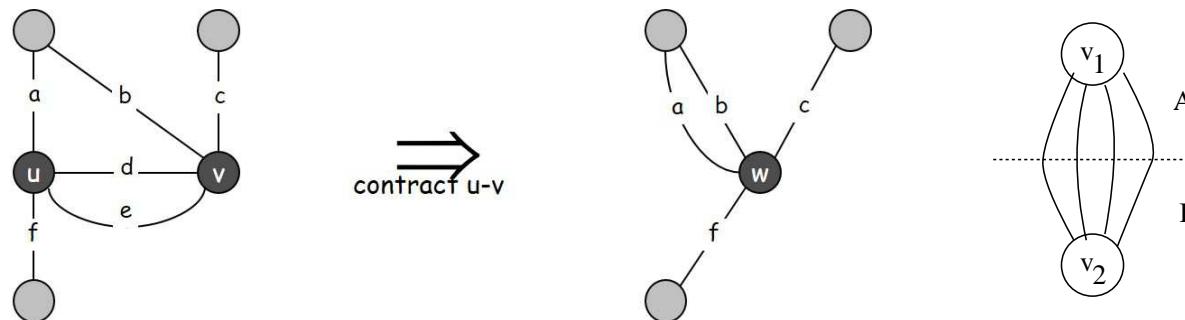
Instance, a connected graph $G(V, E)$.

Goal, find a cut (A, B) (a set of edges whose removal partitions G into two separated components A and B) of minimum cardinality.

- **Applications:** Partition items in a database, identify clusters of related documents, network reliability, network design, circuit design, TSP solvers, etc.
- **Network flow based solution:** Replace every edge $\{u, v\}$ with two arcs $(u, v), (v, u)$; fixed a node s and computes min $s - t$ cut for every $t \in V(G) \setminus \{s\}$.
- **Flow based solution uses $O(n)$ calls to a min $s - t$ cut solver.**
An early intuition is that global min-cut is harder than min $s - t$ cut.
This intuition is not true due to a randomized algorithm [Karger 1995].

Contraction algorithm [Karger 1995]

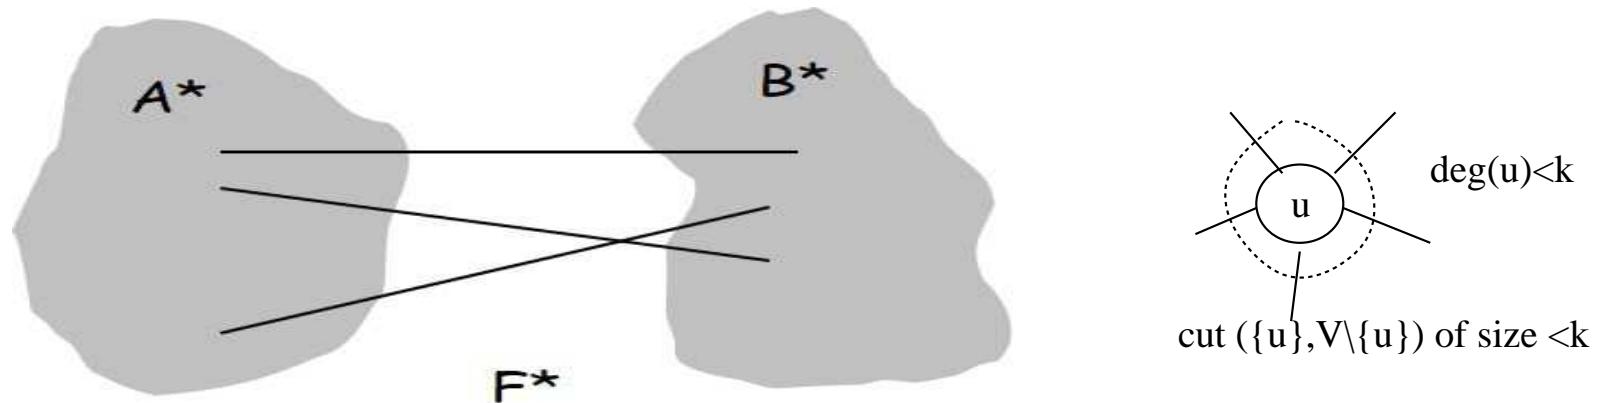
- 1. Select an edge $e = \{u, v\}$ uniformly at random.**
- 2. Contract edge e**
 - replace u and v by a new node w ;
 - for each edge $\{x, y\}$ with $y = u$ or $y = v$, replace y with w ;
 - keep parallel edges and delete self-loops.
- 3. Repeat until graph has two nodes v_1 and v_2 .**
- 4. Return the cut (the set of nodes contracted to form v_1).**



Claim: Given a graph G of n nodes, the contraction algorithm returns a min-cut with probability at least $\frac{2}{n^2}$.

Proof. Consider a global min-cut (A^*, B^*) of G . Let F^* be edge-cut set that gives (A^*, B^*) and $|F^*| = k$.

In the 1st step, an edge in F^* is contracted with probability $k/|E|$. For each node $u \in V$, $\deg(u) \geq k$, as otherwise (A^*, B^*) is not a min-cut, implying $|E| \geq (\frac{1}{2})kn$. Thus, an edge in F^* is contracted in the 1st step with probability at most $\frac{k}{|E|} = \frac{k}{(1/2)kn} = 2/n$.



Let E_j be the event that an edge in F^* is not contracted in iteration j .

$$\begin{aligned}
 & \Pr[E_1 \cap E_2 \cap \dots \cap E_{n-2}] \\
 = & \Pr[E_1] \times \Pr[E_2 | E_1] \times \dots \times \Pr[E_{n-2} | E_1 \cap \dots \cap E_{n-3}] \\
 \geq & \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \left(1 - \frac{2}{n-2}\right) \left(1 - \frac{2}{n-3}\right) \dots \left(1 - \frac{2}{4}\right) \left(1 - \frac{2}{3}\right) \\
 = & \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \left(\frac{n-4}{n-2}\right) \left(\frac{n-5}{n-3}\right) \dots \left(\frac{2}{4}\right) \left(\frac{1}{3}\right) \\
 = & \frac{2}{n(n-1)} \geq \frac{2}{n^2}.
 \end{aligned}$$

□

- **Amplification:** To amplify the probability of finding a min-cut, run the contraction algorithm many times.
- **Claim:** If the contraction algorithm is repeated $n^2 \ln n$ times with independent random edge selections, the probability of failing to find the global min-cut is at most $1/n^2$.

Proof. By independence, the failure probability is at most

$$(1 - \frac{2}{n^2})^{n^2 \ln n} = ((1 - \frac{2}{n^2})^{n^2/2})^{2 \ln n} \leq (e^{-1})^{2 \ln n} = \frac{1}{n^2}.$$

□

- **Running time:** one execution of contraction algorithm takes $O(m)$ time,
 $O(n^2 \ln n)$ runs of the algorithm takes $O(mn^2 \log n)$ time.
- **Improve running time to $O(n^2 \log n)$ [Karger-Stein 1996]:**
 - Early runs of the algorithm are less risky than later runs, the probability of contracting an edge in min-cut hits 0.5 when $n/(\sqrt{2})$ nodes remain.
 - Run contraction algorithm until $n/(\sqrt{2})$ nodes remain to get graph H .
 - Run contraction algorithm twice on H and return the best of two cuts.
- **Best known running time:** $O(m \log^3 n)$ [Karger 2000]; faster than the best known max-flow algorithm or deterministic global min-cut algorithm.

Linearity of Expectation

- The **expectation** of a discrete random variable X taking values v_0, v_1, \dots is defined as:

$$E[X] = \sum_{j=0}^{\infty} v_j \Pr[X = v_j].$$

- Example, waiting for a first success: Coin turns up heads (H) with probability p and tails (T) with probability $1 - p$. How many independent flips X are needed until getting the first H.

$$\begin{aligned} E[X] &= \sum_{j=0}^{\infty} j \Pr[X = j] = \sum_{j=1}^{\infty} j(1-p)^{j-1} p \\ &= \frac{p}{1-p} \sum_{j=1}^{\infty} j(1-p)^j = \left(\frac{p}{1-p}\right)\left(\frac{1-p}{p^2}\right) = \frac{1}{p}. \end{aligned}$$

This result is known as **waiting-time bound**.

Properties of expectation

- If X is a 0/1 random variable, $E[X] = \Pr[X = 1]$.

Proof.

$$E[X] = \sum_{j=0}^{\infty} v_j \Pr[X = j] = \sum_{j=0}^1 j \Pr[X = j] = \Pr[X = 1].$$

□

($\Pr[X = j] = 0$ for $j \geq 2$, $j \Pr[X = j] = 0$ for $j = 0$)

- **Linearity of expectation:** Given two random variables X and Y defined over the same sample space, $E[X + Y] = E[X] + E[Y]$.
- **Application:** Decouple a complex calculation into simpler ones.

Guessing cards

- **Game:** Shuffle n distinct cards, turn them over one at a time, try to guess each card.
- **Memoryless guess:** No psychic abilities, cannot even remember what has been turned over already; guess a card from full deck uniformly at random.
- **Claim:** The expected number of correct guesses is 1.

Proof. Let $X_i = 1$ if the i th guess is correct, otherwise $X_i = 0$.

Let X be the number of correct guesses, that is, $X = X_1 + \dots + X_n$.

$$E[X_i] = \Pr[X_i = 1] = \frac{1}{n}.$$

$$E[X] = \sum_{i=1}^n E[X_i] = \sum_{i=1}^n \frac{1}{n} = 1.$$

□

- **Guess with memory:** Guess a card uniformly at random from cards not seen yet.
- **Claim:** The expected number of correct guesses is $\Theta(\log n)$.

Proof. Let $X_i = 1$ if the i th guess is correct, otherwise $X_i = 0$.

Let X be the number of correct guesses, that is, $X = X_1 + \dots + X_n$.

$$E[X_i] = \Pr[X_i = 1] = \frac{1}{n - (i - 1)}.$$

$$E[X] = \sum_{i=1}^n E[X_i] = \sum_{i=1}^n \frac{1}{n - (i - 1)} = \frac{1}{n} + \dots + \frac{1}{2} + 1 = H(n).$$

$$\ln(n + 1) \leq H(n) \leq 1 + \ln n.$$

□

- **Coupon collector:** Each box of cereal contains a coupon. There are n different types of coupons. Assume all boxes are equally likely to contain each coupon, how many boxes does one need to open before you have at least 1 coupon of each type.
- **Claim:** The expected number of steps is $\Theta(n \log n)$.

Proof. Let phase j be the time between one has j and $j + 1$ distinct coupons. Then $\Pr[\text{get a new coupon at each open in phase } j] = (n - j)/n$.

Let X_j be the number of steps one spends in phase j . Then

$$E[X_j] = 1/p = n/(n - j).$$

Let X be the number of steps in total, i.e. $X = X_0 + X_1 + \dots + X_{n-1}$.

$$E[X] = \sum_{j=0}^{n-1} E[X_j] = \sum_{j=0}^{n-1} \frac{n}{n-j} = n \sum_{i=1}^n \frac{1}{i} = nH(n) = \Theta(n \log n).$$

□

Randomized algorithm for stable matching problem

- Given a set $M = \{m_1, \dots, m_n\}$ employers and a set $W = \{w_1, \dots, w_n\}$ students. each $m \in M$ has a preference list on students defined by bijection $f_m : W \rightarrow \{1, \dots, n\}$, each $w \in W$ has a preference list on employers defined by bijection $g_w : M \rightarrow \{1, \dots, n\}$.
- A matching S between M and W is stable if
 1. S is a perfect matching; and
 2. for any $(m, w'), (m', w) \in S$, $f_m(w') > f_m(w)$ or $g_w(m') > g_w(m)$.
- Problem: find a stable matching S .

Randomized-Stable-Matching

Input: M, W and preference lists for every $w \in W$. **Output:** A stable matching S .

$S = \emptyset$; mark every $m \in M$ and every $w \in W$ unmatched;

while \exists unmatched $m \in M$ **do**

select an unmatched m uniformly at random;

select a $w \in W$ uniformly at random;

if w is not matched then $S = S \cup \{(m, w)\}$ in S ; mark m and w matched;

else

let (m', w) be the pair in S ;

if $g_w(m) > g_w(m')$ then

remove (m', w) from S ; $S = S \cup \{(m, w)\}$; mark m' unmatched;

return S

The algorithm finds a stable matching in $O(n \log n)$ expected running time.

Proof. By a similar analysis for G-S stable matching algorithm, the algorithm finds a stable matching. The expected number of iterations of the while loop is the same as that of collecting all of n types of coupons, which is $O(n \log n)$, □

Maximum 3-Satisfiability

- **Max 3-Sat:** Given a 3-CNF, find an assignment that satisfies as many clauses as possible. The problem is NP-hard.
- **A simple random assignment:** Assign each variable $x_i = 1$ with probability $1/2$ independently.
- **A clause of 3 literals is satisfied with probability $7/8$.**
- **Lemma:** Given a 3-SAT formula of n variables and m clauses, the **expected number of clauses satisfied by a random assignment is $7m/8$.**

Proof. Let Z_j be the random variable with $Z_j = 1$ if clause C_j is satisfied, otherwise $Z_j = 0$. Let Z be the number of clauses satisfied.

$$E[Z] = \sum_{j=1}^m E[Z_j] = 7m/8.$$

□

- **Corollary: For any instance of 3-SAT of n variables and m clauses, there exists a truth assignment that satisfies at least $7m/8$ clauses.**

Proof. Let s_1, \dots, s_r , $r = 2^n$, be the all possible assignments to the n Boolean variables. Let m_i be the number of clauses satisfied by s_i and $m^* = \max_{i=1}^r m_i$. Let Z be the number of clauses satisfied by a randomly selected assignment. Then

$$E[Z] = 7m/8 = \sum_{i=1}^r m_i \Pr[Z = s_i] \leq m^* \sum_{i=1}^r \Pr[Z = s_i] = m^*.$$

Thus, there exists an assignment satisfying at least $7m/8$ clauses. \square

- **Probabilistic method** Prove the existence of a non-obvious property by showing that a random construction produces it with positive probability.

- Is it possible to get a $(7/8)$ -approximation algorithm for the Max 3-Sat problem based the Corollary? It is possible, but there is a hurdle: the probability that a random assignment satisfies $\geq 7m/8$ clauses is not high.
- Lemma: The probability that a random assignment satisfies at least $7m/8$ clauses is at least $1/(8m)$.

Proof. Let p_j be the probability that j clauses are satisfied and $p = \sum_{j \geq 7m/8} p_j$ be the probability that at least $7m/8$ clauses are satisfied. Then

$$\frac{7m}{8} = E[Z] = \sum_{j \geq 0} jp_j = \sum_{j < 7m/8} jp_j + \sum_{j \geq 7m/8} jp_j.$$

If $7m/8$ is an integer, then

$$\begin{aligned}
 \frac{7m}{8} &= \sum_{j \geq 0} jp_j = \sum_{j < 7m/8} jp_j + \sum_{j \geq 7m/8} jp_j \\
 &= \sum_{j \leq 7m/8-1} jp_j + \sum_{j \geq 7m/8} jp_j \\
 &\leq \sum_{j \leq 7m/8-1} (7m/8 - 1)p_j + \sum_{j \geq 7m/8} mp_j \\
 &= \left(\frac{7m}{8} - 1\right) \sum_{j \leq 7m/8-1} p_j + m \sum_{j \geq 7m/8} p_j \leq \left(\frac{7m}{8} - 1\right) \cdot 1 + mp.
 \end{aligned}$$

From this, $mp \geq 1$ and $p \geq 1/m \geq 1/(8m)$.

If $7m/8$ is not an integer, then $\lfloor 7m/8 \rfloor \leq 7m/8 - (1/8)$ and

$$\begin{aligned}
\frac{7m}{8} &= \sum_{j \geq 0} jp_j = \sum_{j < 7m/8} jp_j + \sum_{j \geq 7m/8} jp_j \\
&= \sum_{j \leq \lfloor 7m/8 \rfloor} jp_j + \sum_{j \geq \lceil 7m/8 \rceil} jp_j \\
&\leq \sum_{j \leq (7m/8) - (1/8)} \left(\frac{7m}{8} - \frac{1}{8}\right)p_j + \sum_{j \geq 7m/8} mp_j \\
&= \left(\frac{7m}{8} - \frac{1}{8}\right) \sum_{j < 7m/8} p_j + m \sum_{j \geq 7m/8} p_j \leq \left(\frac{7m}{8} - \frac{1}{8}\right) \cdot 1 + mp.
\end{aligned}$$

From this, $mp \geq (1/8)$ and $p \geq 1/(8m)$. □

- For a 3-CNF of n variables and m clauses, the probability that a random assignment satisfies at least $7m/8$ clauses is at least $1/(8m)$.
- One idea to increase the probability of satisfying $\geq 7m/8$ clauses is to repeatedly generate random assignments until one of them satisfies $\geq 7m/8$ clauses (Johnson's algorithm).
- Theorem: Johnson's algorithm is a $(7/8)$ -approximation algorithm.

Proof. By the previous lemma, one assignment succeeds with probability $\geq 1/(8m)$. By the waiting-time bound, the expected number of trials to find the satisfying assignment is at most $8m$. □

Derandomization: a deterministic $(7/8)$ -approximation algorithm for Max 3-Sat.

- Set each variable x_i of x_1, \dots, x_n to 1 or 0 in the order $1 \leq i \leq n$.
- Let $E[Z|x_1 = b]$, $b \in \{0, 1\}$, be the expected number of satisfied clauses with $x_1 = b$ and every of x_2, \dots, x_n is assigned 1 or 0, each with probability $1/2$ independently.
If $E[Z|x_1 = 1] \geq E[Z|x_1 = 0]$ then set $x_1 = 1$ else set $x_1 = 0$.
- Assume x_1, \dots, x_i have been set ($x_1 = b_1, \dots, x_i = b_i, b_1, \dots, b_i \in \{0, 1\}$). Let $E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = b]$, $b \in \{0, 1\}$, be the expected number of satisfied clauses with $x_{i+1} = b$ and every of x_{i+2}, \dots, x_n is assigned 1 or 0, each with probability $1/2$ independently.

If

$$E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 1] \geq E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 0]$$

then set $x_{i+1} = 1$ **else set** $x_{i+1} = 0$.

Compute $E[Z|x_1 = b_1, \dots, x_i = b_i]$.

$$\begin{aligned} E[Z|x_1 = b_1, \dots, x_i = b_i] &= \sum_{j=1}^m E[Z_j|x_1 = b_1, \dots, x_i = b_i] \\ &= \sum_{j=1}^m \Pr[C_j \text{ satisfied}|x_1 = b_1, \dots, x_i = b_i]. \end{aligned}$$

$\Pr[C_j \text{ satisfied}|x_1 = b_1, \dots, x_i = b_i] = 1$ if C_j is satisfied by one of $x_1 = b_1, \dots, x_i = b_i$, otherwise $\Pr[C_j \text{ satisfied}|x_1 = b_1, \dots, x_i = b_i] = (1 - (\frac{1}{2})^q)$, where q is the number of literals not assigned a value yet.

Example: For $C_j = (x_1 \vee \bar{x}_5 \vee x_7)$,

$$\Pr[C_j \text{ satisfied}|x_1 = 1, x_2 = 0, x_3 = 1] = 1,$$

$$\Pr[C_j \text{ satisfied}|x_1 = 0, x_2 = 0, x_3 = 1] = 1 - (\frac{1}{2})^2 = \frac{3}{4}, \text{ and}$$

$$\Pr[C_j \text{ satisfied}|x_1 = 0, \dots, x_5 = 1, \dots, x_7 = 0] = 0.$$

Given a 3-CNF of n variables and m clauses, the algorithm finds a truth assignment satisfying at least $(7m/8)$ clauses in $O(mn)$ time.

Proof. We show $E[Z|x_1 = b_1, \dots, x_n = b_n] \geq E[Z] \geq (7m/8)$ by induction. For $i = 1$,

$$\begin{aligned} E[Z] &= E[Z|x_1 = 1] \Pr[x_1 = 1] + E[Z|x_1 = 0] \Pr[x_1 = 0] \\ &= \frac{1}{2}(E[Z|x_1 = 1] + E[Z|x_1 = 0]). \end{aligned}$$

If $E[Z|x_1 = 1] \geq E[Z|x_1 = 0]$ then the algorithm sets $x_1 = 1$ and $E[Z|x_1 = 1] \geq E[Z] \geq (7m/8)$.

If $E[Z|x_1 = 0] > E[Z|x_1 = 1]$ then the algorithm sets $x_1 = 0$ and $E[Z|x_1 = 0] \geq E[Z] \geq (7m/8)$.

Assume the statement is true for $i \geq 1$ ($E[Z|x_1 = b_1, \dots, x_i = b_i] \geq E[Z] \geq (7m/8)$), we prove it for $i + 1$.

$$\begin{aligned} E[Z|x_1 = b_1, \dots, x_i = b_i] &= E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 1] \Pr[x_{i+1} = 1] \\ &\quad + E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 0] \Pr[x_{i+1} = 0] \\ &= \frac{1}{2} E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 1] \\ &\quad + \frac{1}{2} E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 0]. \end{aligned}$$

If $E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 1] \geq E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 0]$ then the algorithm sets $x_{i+1} = 1$ and

$$E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 1] \geq E[Z|x_1, \dots, x_i = b_i] \geq E[Z] \geq (7m/8).$$

If $E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 0] > E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 1]$ then the algorithm sets $x_{i+1} = 0$ and

$$E[Z|x_1 = b_1, \dots, x_i = b_i, x_{i+1} = 0] \geq E[Z|x_1, \dots, x_i = b_i] \geq E[Z] \geq (7m/8).$$

It takes $O(m)$ time to assign a value to each x_i . The total running time is $O(mn)$, \square

Two properties of the randomized Max 3-SAT algorithm

- Each variable is set independently
- Conditional expectations can be computed in polynomial time

The derandomization of the randomized Max 3-SAT algorithm is an example of **method of conditional expectations** to derandomize random algorithms with the two properties.

- **Extensions of Max 3-SAT**
 - MAX-SAT: allow one, two, or more literals per clause.
 - Weighted MAX-SAT: each clause is given a non-negative weight, find an assignment to maximize the sum of the weights of the satisfied clauses.
- **Theorem [Asano-Williamson 2000]:** There exists a 0.784-approximation algorithm for MAX-SAT.
- **Theorem [Karloff-Zwick 1997,Zwick+computer 2002]:** There exists a $(7/8)$ -approximation algorithm for the MAX-3SAT with each clause having **at most** 3 literals.
- **Theorem [Håstad 1997]:** Unless P=NP, no ρ -approximation algorithm for MAX-3SAT (hence MAX-SAT) for any $\rho > 7/8$.

A randomized algorithm for k -SAT [Schöning 1999]

Random- k -SAT

Input a k -CNF Φ of n variables and m clauses;

Let $f : \{x_1, \dots, x_n\} \rightarrow \{0, 1\}$ be a randomly selected truth assignment;

Repeat $O(n)$ times

If Φ is satisfied by f then output f and terminate

else

Let C be an unsatisfied clause;

Select one literal l of C uniformly at random;

Update f by flipping the value of l

Algorithm Random- k -SAT, given a satisfiable k -CNF Φ of n variables, finds a satisfying assignment for Φ with probability at least $(\frac{k}{2(k-1)})^n$.

Proof. Let f^* be a satisfying assignment for Φ and f be a randomly selected assignment. Let $d_H(f, f^*) = |\{i|f(x_i) \neq f^*(x_i)\}|$ be the Hamming distance between f and f^* . For each $0 \leq d \leq n$, $\Pr[d_H(f, f^*) = d] = \binom{n}{d}(\frac{1}{2})^n$. Start from f with $d_H(f, f^*) = d$, after one update, $\Pr[d_H(f, f^*) = d - 1] \geq 1/k$, $\Pr[d_H(f, f^*) = d + 1] \leq 1 - (1/k)$, and by a Markov chain analysis, the probability that the algorithm finds a satisfying assignment is at least $(\frac{1}{k-1})^d$. The probability that the algorithm finds a satisfying assignment is

$$\sum_{d=0}^n \binom{n}{d} (\frac{1}{2})^n (\frac{1}{k-1})^d = (\frac{k}{2(k-1)})^n.$$

□

An $O^*((\frac{2(k-1)}{k})^n)$ time randomized algorithm for k -SAT ($O^*((\frac{4}{3})^n)$ time for 3-SAT) can be obtained by repeating Algorithm Random- k -SAT. There is an $O^*((\frac{2(k-1)}{k} + \epsilon)^n)$ time deterministic algorithm for k -SAT ($O^*(1.334^n)$ time for 3-SAT) [Moser and Scheder 2010]. O^* notation hides the polynomial factor $\text{Poly}(n, m)$.

Monte Carlo and Las Vegas Algorithms

- Monte Carlo algorithm: Guaranteed to run in a specific time (e.g., poly-time), likely to find a correct answer (not guaranteed).

Example: Contraction algorithm for global min-cut, Random- k -SAT for k -SAT.

- Las Vegas algorithm: Guaranteed to find a correct answer, likely to run in a specific time (e.g., poly-time) (not guaranteed).

Example: Johnson's MAX-3-SAT algorithm.

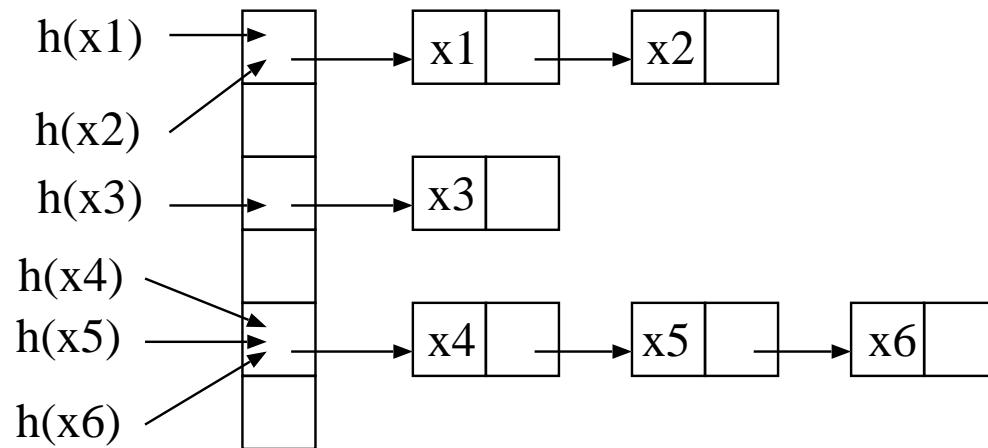
- It is always possible to convert a Las Vegas algorithm into a Monte Carlo algorithm (simply stop the algorithm at some point). No general approach is known for converting a Monte Carlo algorithm into a Las Vegas one.

- **One-sided error:**
 - If the correct answer is NO, always return NO.
 - If the correct answer is YES, return YES with probability at least $1/2$.
- Problem class RP [Monte Carlo]: Set of decision problems solvable with one-sided error in poly-time.
- Problem class ZPP [Las Vegas]: Set of decision problems solvable in **expected** poly-time.
- Theorem: $P \subseteq ZPP \subseteq RP \subseteq NP$.
- Fundamental open problems: Does $P = ZPP$? Does $ZPP = RP$? Does $RP = NP$?

Universal Hashing

- **Dictionary:** Given a set U of possible elements, maintain a subset $S \subseteq U$ so that **inserting**, **deleting** and **searching** in S is efficient.
- **Naive approach:** Create a table T of size $|U|$ and access $T[x]$ when processing $x \in U$. **Challenge:** U can be very large and a table of size $|U|$ is infeasible.
- **Hashing:**
 - Create a smaller table, put some elements of U in the table.
 - To process an $x \in U$, not use x itself, but a function $h(x)$ computed from x to access the table.
 - The function is called **hash function**, the table is called **hash table**, $h(x)$ is called **hash value** of x .
 - Example: create an array A of size n and a hash function $h : U \rightarrow \{0, 1, \dots, n - 1\}$. When processing element x , access array element $A[h(x)]$.

- **Collision in hashing:** $h(x) = h(y)$ but $x \neq y$. A collision is expected after $\Theta(\sqrt{n})$ random insertions of elements into a hash table of size n .
- **Separate chaining:** $A[i]$ stores linked list of elements x with $h(x) = i$.
- Time of accessing an element depends on the length of chains.



- **Good hash functions:** Map m elements uniformly at random to n hash slots.
Average length of chain = m/n . Choose $m \approx n$, $O(1)$ expect access time.
- It is difficult to use a fixed hash function to achieve $O(1)$ access time without the info of the elements (keys) distribution.
- **Randomized approach:** Create a set H of hash functions and select a hash function $h \in H$ at random.
- A **universal family of hash functions** is a set H of hash functions: $U \rightarrow \{0, 1, \dots, n - 1\}$ s.t.
 - for any $x \neq y$, $\Pr_{h \in H}[h(x) = h(y)] \leq 1/n$;
 - selection $h \in H$ at random can be done efficiently; and
 - $h(x)$ can be computed efficiently.

$h_i : U \rightarrow \{0,1\}$, $n=2$

	a	b	c	d	e	f
$h_1(x)$	0	1	0	1	0	1
$h_2(x)$	0	0	0	1	1	1

$$H = \{h_1, h_2\}$$

$$\Pr_{h \in H} [h(a) = h(b)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(c)] = 1$$

$$\Pr_{h \in H} [h(a) = h(d)] = 0$$

not universal

...

	a	b	c	d	e	f
$h_1(x)$	0	1	0	1	0	1
$h_2(x)$	0	0	0	1	1	1
$h_3(x)$	0	0	1	0	1	1
$h_4(x)$	1	0	0	1	1	0

$$H = \{h_1, h_2, h_3, h_4\}$$

$$\Pr_{h \in H} [h(a) = h(b)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(c)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(d)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(e)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(f)] = 0$$

universal

...

- **Lemma:** Assume a hash function h is selected at random from a universal family H of hash functions and used to hash m keys into a table T of size n . For a key x not in S , the expected length of the list that x hashed to is at most m/n .

Proof. Let S be the set of keys in T . For any key $s \in S$, define random variable $X_s = 1$ if $h(x) = h(s)$, otherwise 0. Let X be a random variable counting the total number of collisions with any key in S .

$$\begin{aligned} E_{h \in H}[X] &= E\left[\sum_{s \in S} X_s\right] = \sum_{s \in S} E[X_s] = \sum_{s \in S} \Pr[X_s = 1] \\ &\leq \sum_{s \in S} \frac{1}{n} = \frac{|S|}{n} \leq \frac{m}{n}. \end{aligned}$$

□

- **Corollary:** Using a universal hashing and collision resolution by chaining in a table of size n , it takes expected $\Theta(m)$ time to handle any sequence of m table operations containing $O(n)$ insertions.

- **Constructing a universal family of hash functions**
 - Let p be a prime number. Identify each element $u \in U$ by a key $x = (x_{r-1}, \dots, x_0)$ of r -digit base- p integer ($x = \sum_{i=0}^{r-1} x_i p^i, 0 \leq x_i < p$).
 - Example:** for $p = 2$ and $x = (x_3, x_2, x_1, x_0) = (1, 0, 1, 1)$,
 $x = 1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 = 11$.
 - **Hash function.** Let A be the set of all r -digit base- p integers. For each $a = (a_{r-1}, \dots, a_0) \in A$, define
- $$h_a(x) = \left(\sum_{i=0}^{r-1} a_i x_i \right) \bmod p.$$
- Example:** for $p = 3$, $a = (2, 0, 2, 1)$ and $x = (1, 2, 0, 1)$,
 $h_a(x) = (2 \cdot 1 + 0 \cdot 2 + 2 \cdot 0 + 1 \cdot 1) \bmod 3 = 3 \bmod 3 = 0$.
- **Hash function family:** $H = \{h_a : a \in A\}$ ($h_a : U \rightarrow \{0, 1, \dots, p-1\}$).

Theorem: $H = \{h_a : a \in A\}$ is a universal family of hash functions.

Proof. For two distinct keys $x = (x_{r-1}, \dots, x_0)$ and $y = (y_{r-1}, \dots, y_0)$, we show that $\Pr[h_a(x) = h_a(y)] \leq 1/p (\frac{|\{h_a | h_a(x) = h_a(y)\}|}{|H|} \leq \frac{1}{p})$.

Since $x \neq y$, $x_j \neq y_j$ for some j . Assume $h_a(x) = h_a(y)$. Then

$$h_a(x) - h_a(y) = \sum_{i=0}^{r-1} a_i x_i - \sum_{i=0}^{r-1} a_i y_i = \sum_{i=0}^{r-1} a_i (x_i - y_i) = 0.$$

Thus, $h_a(x) = h_a(y)$ iff

$$a_j (y_j - x_j) \equiv \sum_{i \neq j} a_i (x_i - y_i) \pmod{p}.$$

Let $z = (y_j - x_j)$ and $q = \sum_{i \neq j} a_i (x_i - y_i)$. The equation becomes

$a_j z \equiv q \pmod{p}$ with a_j randomly selected. We show that $a_j z \equiv q \pmod{p}$ has at most one solution $0 \leq a_j < p$, and thus $\Pr[h_a(x) = h_a(y)] \leq 1/p$.

Assume there are two different solutions a_j and a'_j . Then $(a_j - a'_j)z \equiv 0 \pmod{p}$.

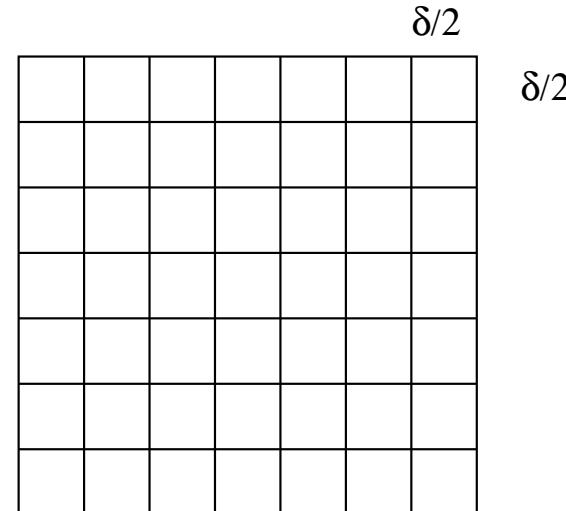
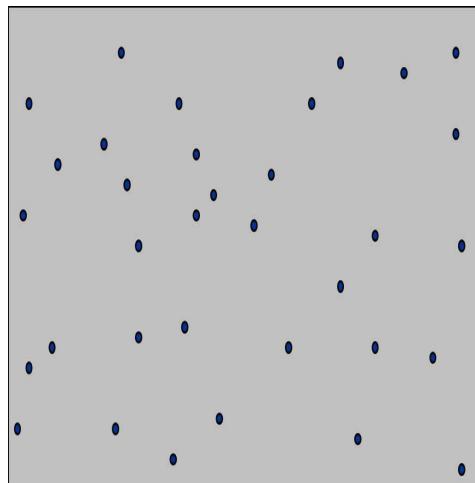
From $z \not\equiv 0 \pmod{p}$, $(a_j - a'_j) \not\equiv 1 \pmod{p}$ and z is not divisible by p . Hence, $(a_j - a'_j)$ is divisible by p . Since p is prime, $(a_j - a'_j) = 0$. □

Summary of universal hash

- Given a set U of elements, maintain a subset $S \subseteq U$ of size m in a table so that each table operation is efficient.
- Select prime p with $m \leq p < 2m$ and create $H = \{h_a, a \in A\}$.
- Space used: $\Theta(m)$.
- Expected number of collisions per table operation is at most 1, implying $O(1)$ expected time per table operation.

Closest Pair Problem

- Given n points $p_1, \dots, p_n, p_i = (x_i, y_i)$ in the 2-dimensional Euclidean plane, find a pair of points with the smallest Euclidean distance $\delta = \min_{1 \leq i < j \leq n} d(p_i, p_j)$ between them.
- Assume all points are in the unit square $P: 0 \leq x_i, y_i \leq 1$.
- Partition the unit square P into $N^2 (\delta/2) \times (\delta/2)$ subsquares $P[s, t]$, $1 \leq s, t \leq N$, for some $0 < \delta < 1$, $N = \lceil 1/(\delta/2) \rceil$.



- If two points p_i, p_j are in a same subsquare, then $d(p_i, p_j) < \delta$.

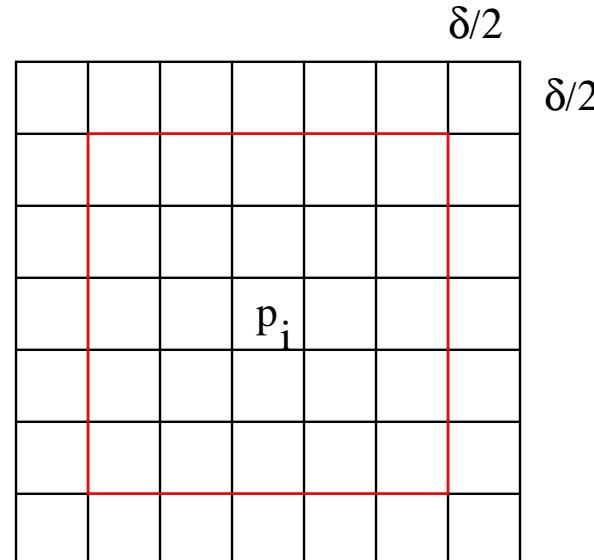
Proof: $d(p_i, p_j) \leq \sqrt{(\delta/2)^2 + (\delta/2)^2} = \delta/\sqrt{2} < \delta$.

- Let $p_i = (x_i, y_i)$ be in subsquare $P[s, t]$. If subsquare $P[s', t']$ has p_j with $d(p_i, p_j) \leq \delta$, then $|s - s'| \leq 2$ and $|t - t'| \leq 2$.

Proof: Assume that $|s - s'| > 2$. Then $|y_i - y_j| > 2(\delta/2) = \delta$. Similarly, $|t - t'| > 2$ implies $|x_i - x_j| > \delta$. Since

$$d(p_i, p_j) \geq \max\{|x_i - x_j|, |y_i - y_j|\}, d(p_i, p_j) > \delta.$$

Let $S(s, t) = \{P[s', t'] || s - s' \leq 2, t - t' \leq 2\}$. Then $|S(s, t)| \leq 25$.



Randomized algorithm for closest pair problem

Order the n points in a random sequence p_1, \dots, p_n ; $\delta = d(p_1, p_2)$;

MakeDictionary to store a subset of $(\delta/2) \times (\delta/2)$ subsquares $P[s, t]$;

/* Create hash table to store subsquares containing a point processed so far

for $i = 1, \dots, n$ do

find the subsquare $P[s, t]$ containing p_i ;

Let $\delta' = \min_{p_j \in P[s', t']} \text{in } S(s, t), j < i d(p_i, p_j)$; /* $P[s', t']$ in dictionary has a p_j

if $\delta' < \delta$ then

delete the current dictionary;

MakeDictionary to store a subset of $(\delta'/2) \times (\delta'/2)$ subsquares $P[s, t]$;

$\delta = \delta'$;

for $j = 1, \dots, i$ do

find the subsquare $P[s, t]$ containing p_j ;

insert $P[s, t]$ into the new dictionary; /* tagged with j

endfor

else insert $P[s, t]$ into the current dictionary; /* tagged with i

endfor

Analysis of the algorithm

- The algorithm correctly keeps the closest pair of p_1, \dots, p_i for every $1 \leq i \leq n$. It has $O(n)$ distance computation, $O(n)$ operations to find subsquare $P[s, t]$, and $O(n)$ MakeDictionary operations.
- Let X_i be a random variable s.t. $X_i = 1$ if $\delta' < \delta$ when p_i is processed, $X_i = 0$ otherwise. The total number of insertions is $n + \sum_{1 \leq i \leq n} iX_i$.
- $\Pr[X_i = 1] \leq 2/i$.

Proof: Let p, q be the two points in p_1, \dots, p_i s.t.

$d(p, q) = \min_{1 \leq i' \neq j \leq i} d(p_{i'}, p_j)$. Since $\delta' < \delta$, either $p = p_i$ or $q = p_i$. As p_1, \dots, p_i are in a random sequence,

$$\Pr[p = p_i \text{ or } q = p_i] \leq (1/i) + (1/i) = 2/i.$$

- Expected number of Insertion operations:

$$E(X) = n + \sum_{1 \leq i \leq n} iE(X_i) \leq n + 2n = 3n.$$

- The algorithm has $O(n)$ expected time plus $O(n)$ expected dictionary operations.
- In expectation, the algorithm uses $O(n)$ hash-function computations and $O(n)$ additional time for finding the closest pair of points.

Proof idea: Applying a universal hashing scheme, the expected total time for lookup the hash table is $O(n)$ (details omitted).

Cache Maintenance Problem

- An algorithm S requests a sequence of data blocks d_1, \dots, d_m , data blocks may not be distinct. A cache can keep at most $k < m$ data blocks. For $1 \leq i \leq m$, when the algorithm requests block d_i ,
 - if d_i is in cache, S accesses d_i in cache (cache hit);
 - if d_i is not in cache and cache has less than k data blocks (cache not full), load d_i into cache and S accesses d_i in cache (cache missing);
 - if d_i is not in cache and cache has k blocks (cache full), evict one block from cache and load d_i into cache, and A accesses d_i in cache (cache missing).
- Optimization goal: minimize the total number of evictions.

A greedy algorithm (Furthest-in-future)

- When a request to a block d not in cache and cache is full, evict a block in cache that comes furthest in the future in the sequence.
- Example, assume cache has size $k = 6$ and the request sequence is:

$a, b, c, d, e, f, g, a, b, c, e, d, a, b, b, a, c, d, e, a, f, a, d, e, f, g, h$

After the 1st six requests, cache holds a, b, c, d, e, f and is full.

The next request is to block g , block f in cache comes furthest in the future in the sequence, so evict f and put g into cache, and cache holds a, b, c, d, e, g right after the request to g .

Furthest-in-future algorithm S_{FF}

- Let d_1, \dots, d_n be the request sequence. For each d_j , let n_j be the smallest index such that $j < n_j$ and $d_j = d_{n_j}$ (n_j is defined $+\infty$ if $d_j \neq d_l$ for every $l > j$).
- $C = \emptyset$. /* Initially, cache C is empty
- For** $i = 1$ **to** n
- if** $d_i = d_j$ **for some** $d_j \in C$ **then change the index** j **to** i
 - else if** $|C| < k$ **then** $C = C \cup \{d_i\}$
 - else** $\{d_{j^*} = \arg \max_{d_j \in C} n_j; C = (C \setminus \{d_{j^*}\}) \cup \{d_i\}\};$
 - endif**
 - endif**
- endfor**

Blocks	a	b	a	b	c	b	c	a	a	b
Cache	a	a	a	a	c	c	c	a	a	a
k=2		b	b	b	b	b	b	b	b	b

Theorem FF: Algorithm furthest-in-future S_{FF} has the minimum number of evictions.

Proof. For a sequence $\sigma = d_1, \dots, d_m$ of blocks and $1 \leq i \leq m$, let σ_i be the subsequence of the first i blocks d_1, \dots, d_i . For a schedule S , let $f_S(\sigma_i)$ be the number of evictions by S on σ_i , $f_S(\sigma_m) = f_S(\sigma)$.

A schedule is **reduced** if it only puts a block in cache when the block is requested, otherwise **unreduced**.

When a data block is put into cache at the time it is not requested, we count this as an **unreduced** event.

The proof is divided into two parts. Part 1: for any unreduced schedule S , there is a reduced schedule S' s.t. $f_{S'}(\sigma) \leq f_s(\sigma)$. Part 2: For any reduced schedule S , $f_{S_{FF}}(\sigma) \leq f_S(\sigma)$.

- Proof for Part 1

Claim: Any unreduced schedule S can be transformed to a reduced schedule S' s.t.
 $f_{S'}(\sigma) \leq f_S(\sigma)$.

We prove the claim by induction on the number i of unreduced events.

For $i = 0$, claim holds. Assume claim true for $0 \leq i - 1$ and prove it for i . Let d be the block put in cache at the time t in the i th unreduced event. Then d is either evicted at time $t' > t$ before the next request to d or d is requested at time $t' > t$ before d is evicted. We simply transform S to S' by removing the operation to put d into cache at time t .

- Proof for Part 2

Claim: For any reduced schedule S , $f_S(\sigma) \geq f_{S_{FF}}(\sigma)$.

We prove the following statement by induction: Let S be any reduced schedule that $f_S(\sigma_i) \geq f_{S_{FF}}(\sigma_i)$ for $i \geq 0$. Then there is a reduced schedule S' that S and S' have same cache contents and $f_S(\sigma_j) \geq f_{S'}(\sigma_j) \geq f_{S_{FF}}(\sigma_j)$ for some j with $i + 1 \leq j \leq m$.

For $i = 0$ and any schedule S , $f_S(\sigma_i) = f_{S_{FF}}(\sigma_i)$. When access block $d_{i+1} = d_1$, let S' follow what S does. Then S and S' have same cache contents and $f_S(\sigma_{i+1}) = f_{S'}(\sigma_{i+1}) \geq f_{S_{FF}}(\sigma_{i+1}) = 0$. The statement is true and we get the induction base.

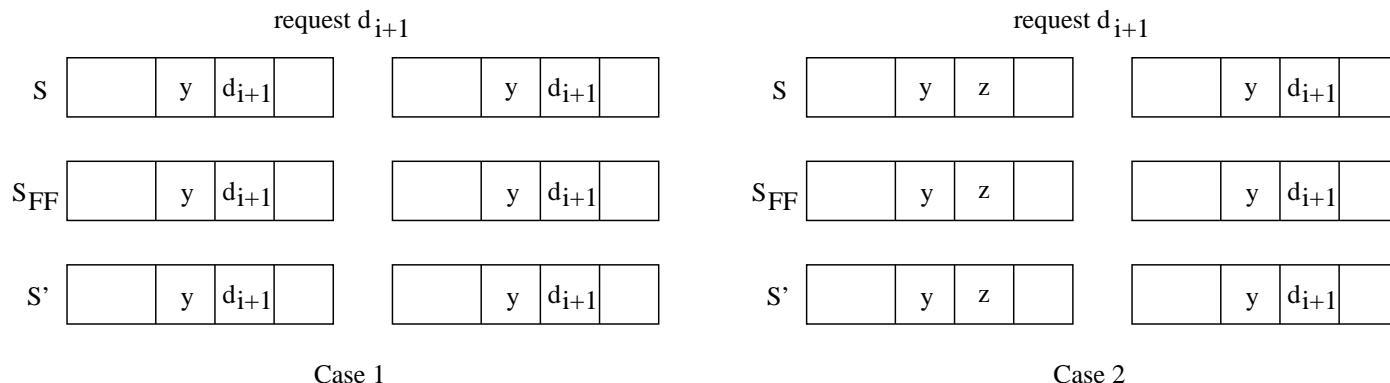
(Proof for Part 2 continued)

Assume the statement holds for $i \geq 0$ and prove it for $i + 1$.

When access d_{i+1} , S and S' have same cache contents. There are 3 cases:

1. d_{i+1} is in cache and no eviction by S ;
2. d_{i+1} is not in cache, S and S_{FF} evict a same block from cache;
3. d_{i+1} is not in cache, S evicts block z and S_{FF} evicts $y \neq z$.

For Cases 1 and 2, let S' follow S . Then S and S' have same cache contents and $f_S(\sigma_{i+1}) = f_{S'}(\sigma_{i+1}) \geq f_{S_{FF}}(\sigma_{i+1})$.



(Proof for Part 2 continued)

For Case 3, first let S' follow S_{FF} to evict y . Let j be the minimum index s.t. $i + 2 \leq j$ and one of the following events happens:

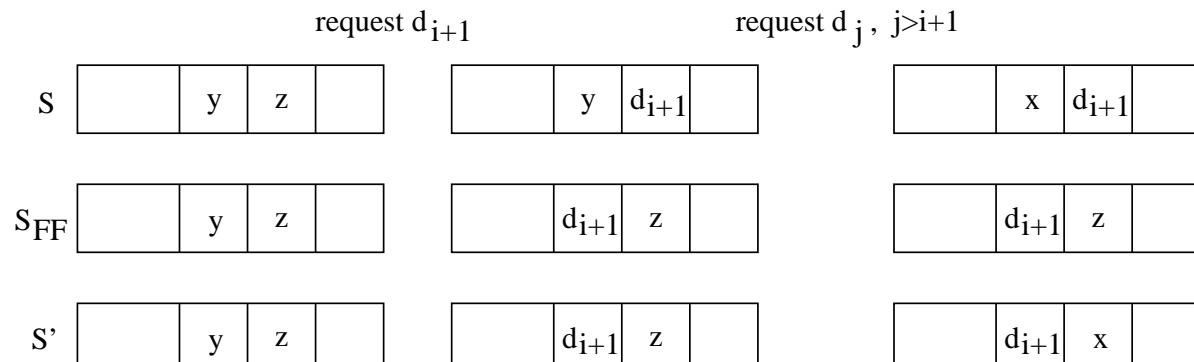
- (A) before z is put back to cache, S evicts y for a request to $x \neq y$;
- (B) S evicts some w for a request to z .

Then S' follows S on d_{i+2}, \dots, d_{j-1} .

When a request for d_j :

- (A) $d_j = x \neq y$, z not in cache, and S evicts y to put x in cache.

In this case, S' does not have x in cache. So, S' evicts z to put x in cache. After this, S' and S have same cache contents and $f_S(\sigma_j) = f_{S'}(\sigma_j) \geq f_{S_{FF}}(\sigma_j)$.



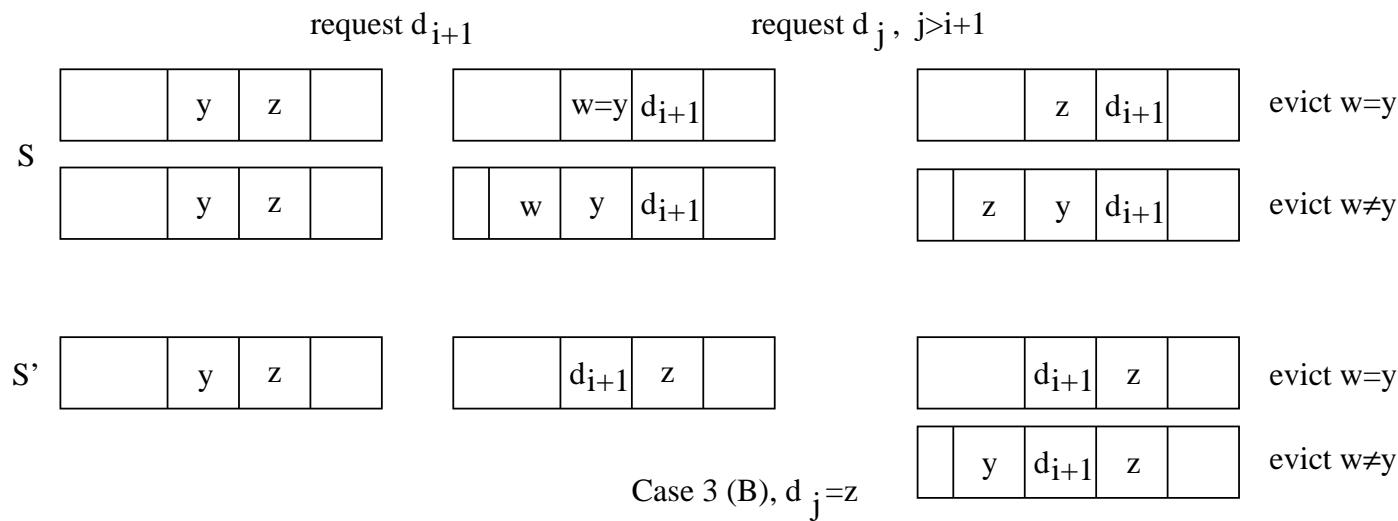
Case 3 (A), $d_j = x \neq y$

(B) $d_j = z$ and S evicts w .

If $w = y$, then S' does not evict any block. After this, S and S' have same cache contents and $f_S(\sigma_j) > f_{S'}(\sigma_j) \geq f_{S_{FF}}(\sigma_j)$. If $w \neq y$, then S' evicts w and put y into cache. After this, S and S' have same cache contents. S' becomes an unreduced schedule. By the proof for Part 1, we can transform S' into reduced. Then $f_S(\sigma_j) \geq f_{S'}(\sigma_j) \geq f_{S_{FF}}(\sigma_j)$.

Summarize above, the statement is proved.

From the proofs for Part 1 and Part2, the theorem holds.



An online algorithm

- When algorithm A requests block d_i , no knowledge on future blocks

d_{i+1}, \dots, d_m .

- Least-recently-used (LRU) algorithm

For $k < i \leq m$, let d_{j_1}, \dots, d_{j_k} be the blocks in cache and t_1, \dots, t_p be the sequence of accesses to the blocks in cache. For each d_{j_l} in cache, let q_{j_l} be the maximum index s.t. $q_{j_l} \leq p$ and d_{j_l} is accessed in $t_{q_{j_l}}$.

For $1 \leq i \leq m$ do

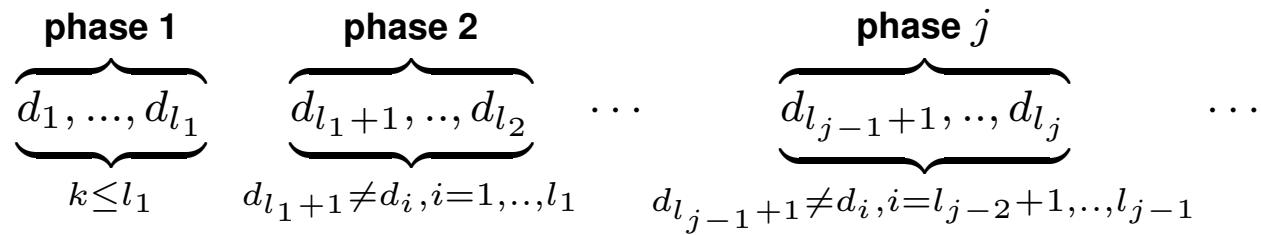
if cache is full then let $j^* = \arg \min_{q_{j_l}} d_{j_l}$, evict d_{j^*} and load d_i to cache.

- LRU algorithm has at most $k \cdot \text{opt} + k$ evictions, where opt is the value of optimal solution (proof later).

Blocks	a	b	a	b	c	b	c	a	a	b
Cache	a	a	a	a	c	c	c	c	c	b
$k=2$		b	b	b	b	b	b	a	a	a

A general template, marking algorithms

- Consider requests to data blocks d_1, d_2, \dots, d_m as phases of requests, each phase requested k distinct blocks, a new phase starts when a block different from the k blocks is requested.



- At start of a phase, every block is unmarked (not recently requested).
When a block d is requested in the phase, d is marked (recently requested).
When an eviction is needed, evict an unmarked block.
When all blocks in cache become marked and an eviction is needed, finish the current phase and start a new phase.

- **Marking algorithms, evict a block not recently requested.**

Initially, every block not requested is considered not recently requested

$i = 1;$

while $i \leq m$ **do**

unmark all blocks in cache C ; /*start of a phase,

mark every block in C **not recently requested**

while ($C \neq$ full **or** $\exists d \in C$ unmarked) **and** $i \leq m$ **do**

mark d_i ; /***mark** d_i recently requested

if $C \neq$ full **and** $d_i \notin C$ **then load** d_i **into** C

else if ($C =$ full **and** $\exists d \in C$ unmarked) **and** $d_i \notin C$ **then**

evict an unmarked block; load d_i **into** C ; $i = i + 1$

endif

endif

endwhile /* end of a phase

endwhile

- **LRU algorithm is a marking algorithm**

Analysis of marking algorithm

- In each phase, there are at most k evictions.

Proof: After k evictions, every block in C becomes marked.

- If there are r phases, there are at most kr evictions.
- An optimal algorithm has at least $r - 1$ evictions.

Proof: Let S_i be the subsequence of blocks processed in the i th phase. Then for each S_i with $1 \leq i \leq r - 1$, S_i has $k + 1$ distinct blocks and an optimal algorithm must make at least one eviction.

- For any marking algorithm, the number of evictions is at most $k \cdot \text{opt} + k$.

Proof: $kr = k(r - 1) + k \leq k \cdot \text{opt} + k$.

- An example: assume $m/(k + 1) = l$ is an integer and $d_1, \dots, d_m = S_1, S_2, \dots, S_l$, where each $S_i = d_1, \dots, d_k, d_{k+1}$. Then $\text{opt} = m/k$ and the number of evictions by LRU is $m - k = k(m/k) - k$.

Randomized marking algorithm

- Initially, every block not requested is considered not recently requested

$i = 1;$

while $i \leq m$ **do**

unmark all blocks in cache C ; /*start of a phase,

mark every block in C **not recently requested**

while ($C \neq$ full **or** $\exists d \in C$ unmarked) **and** $i \leq m$ **do**

mark d_i ; /*mark d_i recently requested

if $C \neq$ full **and** $d_i \notin C$ **then load** d_i **into** C

else if ($C =$ full **and** $\exists d \in C$ unmarked) **and** $d_i \notin C$ **then**

evict an unmarked block selected uniformly at random;

load d_i **into** C ; $i = i + 1$

endif

endif

endwhile /* end of a phase

endwhile

Analysis of randomized marking algorithm

- A block d requested in phase j is called **fresh** if d is not marked (requested) in phase $j - 1$, otherwise called **stale**. Let c_j be the number of fresh blocks in phase j . Let r be the number of phases. Let $f^*(\sigma)$ be the number of evictions by an optimal schedule.
- **Claim 1:** $f^*(\sigma) \geq (1/2) \sum_{j=2}^r c_j$.

Proof. Let $f_j^*(\sigma)$ be the number of evictions in phase j by the optimal schedule. Then $f^*(\sigma) = \sum_{j=1}^r f_j^*(\sigma)$. In any phase $j \geq 1$, there are k requests to k distinct blocks and there are requests to c_{j+1} distinct fresh blocks in phase $j + 1$. So there are at least $k + c_{j+1}$ requests to distinct blocks in phases j and $j + 1$. This implies that the optimal schedule has at least c_{j+1} evictions in phases j and $j + 1$. So, $f_j^*(\sigma) + f_{j+1}^*(\sigma) \geq c_{j+1}$. Then,

$$2f^*(\sigma) = 2 \sum_{j=1}^r f_j^*(\sigma) \geq \sum_{j=1}^{r-1} (f_j^*(\sigma) + f_{j+1}^*(\sigma)) \geq \sum_{j=1}^{r-1} c_{j+1}.$$

□

- **Claim 2: Let X be the number of evictions by the randomized algorithm. For every sequence σ , $E[X] \leq H_k \sum_{j=2}^r c_j$, H_k is the harmonic function.**

Proof. There are k requests to unmarked blocks in each phase, c_j blocks are fresh and $k - c_j$ are stale. Let X_j be the number of evictions in phase $j \geq 2$. Then $X_j \geq c_j$ since each fresh block is not marked in phase $j - 1$ and is not in cache when requested in phase j .

For a stale block d , d may be in cache when requested in phase j . If d is in cache then no eviction happens, otherwise an eviction happens. What is the probability an eviction happens?

At the beginning of a phase, all k blocks are stale by the initialization of a phase. For the l th request to a stale block d initially in cache, when d is requested, $l - 1$ initially stale blocks become marked and $k - l + 1$ blocks remain stale.

Let c be the number of requests to fresh blocks before d is requested. Then c initially stale blocks are evicted and not in cache. Each of the $k - l + 1$ has the same probability not in cache. So $\Pr[d \text{ not in cache}] = \frac{c}{k-l+1} \leq \frac{c_j}{k-l+1}$. Then

$$E[X_j] \leq c_j + \sum_{l=1}^{k-c_j} \frac{c_j}{k-l+1} \leq c_j \left(1 + \sum_{l=c_j+1}^k \frac{1}{l}\right) = c_j (1 + H_k - H_{c_j}) \leq c_j H_k.$$

From this,

$$E[X] = \sum_{j=2}^r E[X_j] \leq H_k \sum_{j=2}^r c_j.$$

□

- **Theorem: The expected number of evictions by the randomized marking algorithm is $O((\log k) \cdot f^*(\sigma))$.**

Proof. The theorem follows from Claims 1 and 2. □

Chernoff Bounds

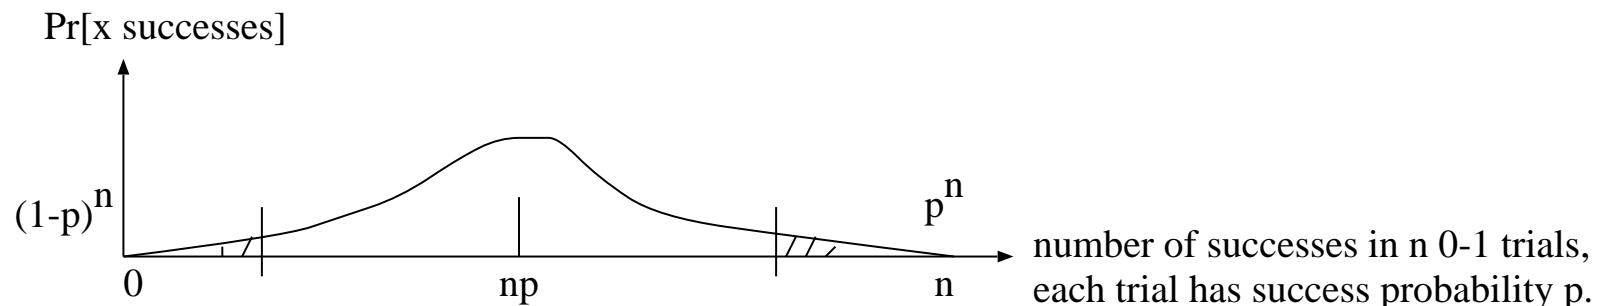
- **Theorem (above mean): Assume X_1, \dots, X_n are independent 0 – 1 random variables. Let $X = X_1 + \dots + X_n$. Then for any $\mu \geq E[X]$ and any $\delta > 0$,**

$$\Pr[X > (1 + \delta)\mu] < \left(\frac{e^\delta}{(1 + \delta)^{1+\delta}} \right)^\mu$$

Intuition: Sum of independent 0 – 1 random variables is tightly centred on the mean.

- **Looser but more convenient bound in practice**

$$\Pr[X > (1 + \delta)\mu] \leq \left(\frac{1}{e} \right)^{\delta^2 \mu / (2 + \delta)}$$



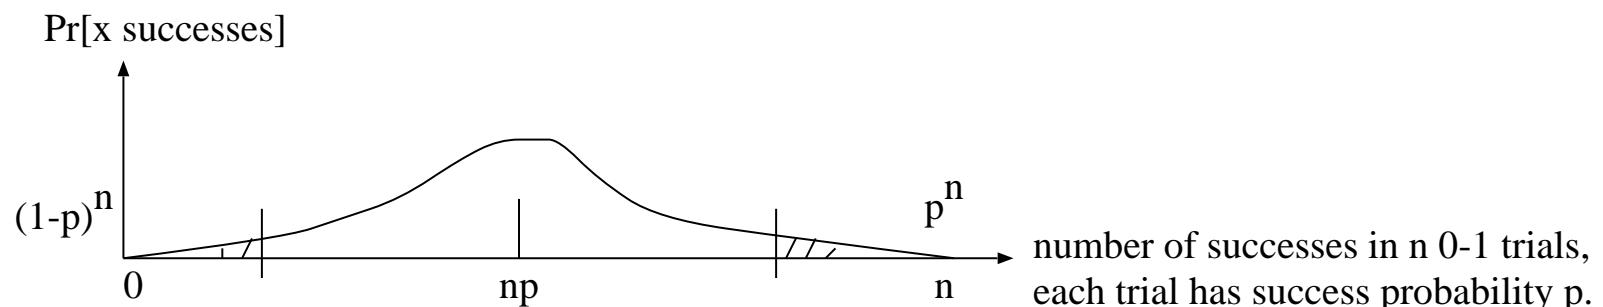
- **Theorem (below mean): Assume X_1, \dots, X_n are independent 0 – 1 random variables. Let $X = X_1 + \dots + X_n$. Then for any $\mu \leq E[X]$ and any $0 < \delta < 1$,**

$$\Pr[X < (1 - \delta)\mu] < \left(\frac{e^{-\delta}}{(1 - \delta)^{1-\delta}} \right)^\mu$$

Intuition: Sum of independent 0 – 1 random variables is tightly centred on the mean.

- **Looser but more convenient bound in practice**

$$\Pr[X < (1 - \delta)\mu] \leq \left(\frac{1}{e} \right)^{\delta^2 \mu / 2}$$



Load Balancing

- Load balancing problem

Instance: System in which m jobs arrive in a stream and need to be processed immediately on n identical processors.

Objective: Find an assignment to balance the workload across processors.

- **Centralized controller:** Assign jobs in round-robin manner. Each processor receives at most $\lceil m/n \rceil$ jobs.
- **Decentralized controller:** Assign jobs to processors uniformly at random.

How likely is it that some processor is assigned "too many" jobs?

- **Analysis** Let X_i be the number of jobs assigned to processor i .

Let $Y_{ij} = 1$ if job j assigned to processor i , otherwise 0.

Then $E[Y_{ij}] = 1/n$, $X_i = \sum_{1 \leq j \leq m} Y_{ij}$ and $\mu = E[X_i] = m/n$.

Applying Chernoff bounds with $\delta = c - 1$ gives

$$\Pr[X_i > c\frac{m}{n}] < \left(\frac{e^{c-1}}{c^c}\right)^{m/n}.$$

Let $\gamma(n)$ be the number x s.t. $x^x = n$, and choose $c = e\gamma(n)$.

$$\begin{aligned} \Pr[X_i > c\frac{m}{n}] &< \frac{e^{(c-1)\frac{m}{n}}}{c^{c\frac{m}{n}}} < \left(\frac{e}{c}\right)^{c\frac{m}{n}} \\ &= \left(\frac{1}{\gamma(n)}\right)^{e\gamma(n)\frac{m}{n}} < \left(\frac{1}{\gamma(n)}\right)^{2\gamma(n)\frac{m}{n}} = \frac{1}{n^{2m/n}}. \end{aligned}$$

Union bound implies that with high probability at least $1 - 1/n$ no processor receives more than $e\gamma(n)\frac{m}{n} = \Theta(\frac{\log n}{\log \log n} \frac{m}{n})$ jobs.

Theorem: Assume the number of jobs is $m = 16n \ln n$ jobs. Each of the n processor handles $\mu = 16 \ln n$ jobs on average. With high probability, each processor will have at least half and at most twice the average load.

Proof. Let X_i and Y_{ij} be defined as in the analysis. Applying Chernoff bounds with $\delta = 1$ gives

$$\Pr[X_i > 2\mu] < \left(\frac{e}{4}\right)^{16 \ln n} < \left(\frac{1}{e^2}\right)^{\ln n} = \frac{1}{n^2}.$$

$$\Pr[X_i < \frac{1}{2}\mu] < e^{-\frac{1}{2}(\frac{1}{2})^2(16 \ln n)} = \frac{1}{n^2}.$$

Union bound implies that each processor has load between half and twice the average with probability at least $1 - 2/n$. □