

Due 23:59 Nov 16 (Sunday). There are 100 points in this assignment.

Submit your answers (**must be typed**) in pdf file to CourSys

<https://coursys.sfu.ca/2025fa-cmpt-705-x1/>.

Submissions received after 23:59 will get penalty of reducing points: 20 and 50 points deductions for submissions received at [00 : 00, 00 : 10] and (00 : 10, 00 : 30] of Nov 17, respectively; no points will be given to submissions after 00 : 30 of Nov 17.

1. (Chapter 11 Problem 1 of the text book) 15 points

There are n containers $1, \dots, n$ of weights w_1, \dots, w_n and some trucks, each truck can hold at most K units of weight (w_i and K , $w_i \leq K$, are positive integers). Multiple containers can be put on a truck subject to the weight restriction K . The problem is to use a minimum number of trucks to carry all containers. A greedy algorithm for this problem is as follows: start with an empty truck and put containers $1, \dots, j$ on the truck to have this truck loaded (i.e., $\sum_{1 \leq i \leq j} w_i \leq K$ and $\sum_{1 \leq i \leq j+1} w_i > K$); then put containers $j+1, j+2, \dots$ on a new empty truck to have this truck loaded; continue this process until all containers are carried.

- (a) Give an example of a set of containers and a value K to show that the greedy algorithm above does not give an optimal solution.

Solution:

Consider this example $k = 5$, w_1 , to w_4 equal to 3, 4, 5, 2 respectively. The greedy algorithm returns 4 trucks, (each container in one truck) but if items with weights 2 and three put in one truck, the total number of trucks are 3. Thus, the greedy algorithm does not always return the optimal solution.

- (b) Prove the greedy algorithm is a 2-approximation algorithm.

Solution:

Let m be the total number of trucks from the Greedy algorithm and the optimal number of trucks be opt . We have:

$$\sum_{i=1}^n w_i > n \times opt$$

let L_j be the load of j th truck from the solution of the Greedy Algorithm. We have:

$$\sum_{j=1}^m L_j = \sum_{i=1}^n w_i$$

Let w_j be the first container in the sequence that could not be loaded into truck j . We have:

$$K < L_j + w_j, \quad \sum_{j=1}^m w_j < \sum_{i=1}^n w_i$$

Sum the first inequality over all m trucks to get:

$$K \times m < \sum_{j=1}^m L_j + \sum_{j=1}^m w_j = \sum_{i=1}^n w_i + \sum_{j=1}^m w_j < \sum_{i=1}^n w_i + \sum_{i=1}^n w_i = 2 \times \sum_{i=1}^n w_i \leq 2k \times opt$$

Thus: $m < 2opt$

2. (Chapter 11 Problem 2 of the text book) 10 points

Given a set S of strings, for any two strings p and q in S , a distance $d(p, q) \geq 0$ is defined. Given a similarity threshold value $\Delta \geq 0$, two strings p and q are called similar if $d(p, q) \leq \Delta$. A subset R of S is called a representative set of S if for any string $p \in S$, there is a string $q \in R$ with $d(p, q) \leq \Delta$. The minimum representative set problem is to find a representative set of minimum size.

Give a polynomial time $O(\log n)$ -approximate algorithm for the problem and analyze the algorithm.

Solution:

Run the following algorithm to find S_q for each q in S . This algorithm runs in $O(n^2)$ time.

Input: S and $d(p, q)$, for each p, q in S .

Output: $S_q = \{p | d(p, q) \leq \Delta\} \cup \{q\}$

for each string q in S :

$S_q = \{q\}$

for each string p in S :

if $d(p, q) \leq \Delta$

$S_q = S_q \cup \{p\}$

End if

End for

End for

S_q , for each $q \in S$ is a set cover for S as S_q contains at least q , thus $\cup_q S_q = S$

Apply the Greedy Set Cover algorithm introduced in lecture notes. This algorithm is a $(1 + \ln n)$ -approximation. Prove is provided in the lecture notes. The first step runs in $O(n^2)$ and the greedy algorithm runs in poly time. Thus, the algorithm runs in poly time.

3. (Chapter 11 Problem 5 of the text book) 15 points

(a) Consider a load balancing problem instance of 10 machines and n jobs $S = \{1, \dots, n\}$ with $1 \leq t_i \leq 50$ for every i and $\sum_{i=1}^n t_i \geq 3000$. Prove that the greedy algorithm discussed in class finds a solution of makespan T with $T \leq (1.2)(\sum_{1 \leq i \leq n} t_i)/10$ for this instance.

Solution:

let L be the solution of the greedy algorithm:

We have:

$$L \geq \frac{\sum t_i}{m} = \frac{3000}{100} = 300$$

Thus, the load on the machine with longest run time is at least 300. Since each job runs in at most 50, thus there are at least 6 jobs on the machine with the longest running time.

Let L_i be the machine with the longest running time and t_l be the last job assigned to this machine. Since t_l is the last assigned job to L_i , it is shorter than all other jobs assigned to L_i and at least 5 other jobs are assigned to L_i . Thus, $\sum_{j=1}^5 t_j \leq L_i$ when t_1 to t_5 be those 5 other jobs. We have:

$$t_l \leq \frac{\sum_{j=1}^5 t_j}{5} \leq \frac{L_i}{5} \leq \frac{\sum t_i}{5m}$$

By the algorithm, at the time this job was assigned machine i has the least load. Thus, $L_i - t_l \leq L_k$, for all machine k . Sum this inequality over all m machines, we have:

$$m \times L_i - m \times t_l \leq \sum_k^m L_k = \sum_i^n t_i$$

thus:

$$l_i \leq \frac{\sum_i^n t_i}{m} + t_l \leq \frac{\sum_i^n t_i}{m} + \frac{\sum_i^n t_i}{5m} = 1.2 \frac{\sum_i^n t_i}{m}$$

Put $m = 10$ and $T = L_i$ we get the requested upper bound.

(b) Implement the greedy algorithm to find a solution for the load balancing problem instance in (a) and compare the solution with the lower bound $(\sum_{1 \leq i \leq n} t_i)/10$ on the makespan of the instance (each t_i can be generated randomly). Report your results for T , $(\sum_{i=1}^n t_i)/10$ and $T/((\sum_{i=1}^n t_i)/10)$ with $n = 100$.

Solutions:

Number of jobs n : 100

Number of machines m : 10

Makespan T : 308

$(\sum_{i=1}^n t_i)/10 : 306.90$

$T/((\sum_{i=1}^n t_i)/10) : 1.0036$

4. (Chapter 11 Problem 7 of text book) 15 points

Given a set of customers $\{1, 2, \dots, n\}$, each customer has a value v_i and is shown to one of the advertisements (ads) A_1, \dots, A_m . A selection of ads to customers is to assign one

ad to each customer. For a set C_j of customers assigned ad A_j , the total value of C_j is $v(C_j) = \sum_{i \in C_j} v_i$. The spread of the selection is $\min_{1 \leq j \leq m} \{v(C_j)\}$. The problem of finding the maximum spread is NP-hard. Give a $(1/2)$ -approximation algorithm for finding the maximum spread for any input instance with $v_i \leq (\sum_{k=1}^n v_k)/(2m)$ for $1 \leq i \leq n$, and analyze the algorithm.

Solution:

Use an algorithm similar to that of load balancing:

LPT – ($m, n, v_1, v_2, \dots, v_n$)

Input: m ads and n customers with value v_i .

Output: A allocation of customers to ads.

```

Sort customers so that  $v_1 \geq v_2 \geq \dots \geq v_n$ ;
for  $j = 1$  to  $m$  do
     $v(C_j) = 0$ ; /* load on machine i */
     $C_j = \emptyset$ ; /* jobs assigned to machine i */
end for
for  $i = 1$  to  $n$  do
     $j = \operatorname{argmin}_{1 \leq k \leq m} \min v(C_k)$ ; /* ad j has the smallest value */
     $C_j = C_j \cup \{i\}$ ; /* assign job j to machine i */
     $v(C_j) = v(C_j) + t_i$ ; /* update load of machine i */
end for
Output  $v(C_j)$  and  $C_j$  for  $1 \leq i \leq m$ 
```

Running time:

Implementation takes $O(n \log n)$ time using a priority queue and merge-sort

Prove:

Let opt be the optimal spread. Thus:

$$\frac{\sum_{i=1}^n v_i}{m} \geq opt$$

V be the solution of the greedy algorithm., and v_l be the last customer assigned to the ad j which has the largest value on the solution of the greedy algorithm.

by definition:

$$v(C_j) \geq \frac{\sum_{k=1}^m v(C_k)}{m} = \frac{\sum_{i=1}^n v_i}{m}$$

At the time v_l assigned to ad j , it has the least value, even less than the spread. Thus,

$$V \geq v(C_j) - v_l \geq \frac{\sum_{i=1}^n v_i}{m} - v_l$$

from the question, we have: $v_i \leq (\sum_{k=1}^n v_k)/(2m)$. Thus:

$$V \geq \frac{\sum_{i=1}^n v_i}{m} - v_l \geq \frac{\sum_{i=1}^n v_i}{m} - \frac{\sum_{i=1}^n v_i}{2m} = \frac{\sum_{i=1}^n v_i}{2m} \geq \frac{opt}{2}$$

Thus, the algorithm is $1/2$ approximate.

5. (Chapter 11 Problem 8 of text book) 15 points

For every instance of the load balancing problem discussed in class, there exists an order of the jobs so that when Greedy-Balance processes the jobs in this order, it produces an assignment of jobs to machines with the minimum possible makespan.

Solution:

Yes. We can construct this order for every input instance.

First, find the minimum possible makespan time allocation. If there are more than one solution with the minimum makespan, pick a solution for which the minimum of run time of all machines is maximized. Let S_j be the jobs assigned to machine j in this optimal solution.

Make sequence R as follow: Put the first jobs in S_1 to S_m as the first m elements of sequence R . The greedy algorithm assigns each job to the associated machine.

After that, if machine j has the lowest run time, put one job from S_j in the sequence. If two machine have the minimum run time, choose a job from S_j with lowest j , as this is the tie breaking rule in the greedy algorithm.

If jobs are fed to the greedy algorithm in order of sequence R , the greedy algorithm returns optimal solution.

Prove:

By contradiction assume we cannot follow the algorithm. That is at the end of the algorithm, machine j has never had the lowest run time, but there is a job a_m in S_j that is not assigned to sequence R .

if $L_j - a_m$ is larger than all run times, then can we move job a_m from machine j and another machine and decrease the makespan, which is a contradiction with the fact that it was optimal.

Let machine k has the smallest run time. thus: $L_k < L_j - a_m$

Hence if we move job a_m from machine j and put it on machine k the run time of machine k will increase. That means L_k could be larger while makespan the same, which is a contradiction with the fact that the machine with solution has the maximum, of minimum run time for all machines.

6. (Chapter 10 Problem 2 of text book) 15 points

Given a 3-SAT instance Φ of n variables, a truth assignment $f : \{x_1, \dots, x_n\} \rightarrow \{0, 1\}$ and an integer $d \geq 0$, the procedure $\text{Explore}(\Phi, f, d)$ below answers whether there is a truth assignment $f' : \{x_1, \dots, x_n\} \rightarrow \{0, 1\}$ such that $d(f, f') = |\{i | f(x_i) \neq f'(x_i)\}| \leq d$ and f' satisfies Φ .

$\text{Explore}(\Phi, f, d)$

```
if  $f$  satisfies  $\Phi$  then return YES
else if  $d = 0$  then return NO
```

else

let $C = (l_1 \vee l_2 \vee l_3)$ be a clause of Φ that is not satisfied by f ;

let f_i ($i = 1, 2, 3$) be the truth assignment obtained from

f by inverting the assigned value to l_i ;

if $\text{Explore}(\Phi, f_1, d - 1) = \text{YES}$ then return YES;

if $\text{Explore}(\Phi, f_2, d - 1) = \text{YES}$ then return YES;

if $\text{Explore}(\Phi, f_3, d - 1) = \text{YES}$ then return YES;

return NO

- (a) Prove $\text{Explore}(\Phi, f, d)$ returns YES iff there is a satisfying assignment f' with $d(f, f') \leq d$. Analyze the running time of $\text{Explore}(\Phi, f, d)$ as a function of n and d .

Solution:

If there is another assignment f' with $d(f, f') \leq d$, then if invert at most d variables in f , it should be satisfiable. at each loop the algorithm changes one of them and try all possible variables in the clauses that are not satisfied. Thus, after at most d loop, it will reach to a satisfiable assignment and return yes

Proof by contradiction. Assume if return YES then there is no satisfying assignment f' with $d(f, f') \leq d$. The algorithm will invert at most d variables and get a satisfiable solution. Thus, there is a satisfiable solution with distance less than d . Which is a contradiction.

Running time:

$$T(n, d) \leq 3 \times T(n, d - 1) + O(m + n)$$

Thus, based on the master theorem the running time is $O(n^{O(1)}3^d)$

- (b) Using $\text{Explore}(\Phi, f, d)$ as a subroutine, give an algorithm which decides whether a 3-SAT instance is satisfiable or not in $O(n^{O(1)}(\sqrt{3})^n)$ time.

The solution should have polynomial many calls to explore with $d = n/2$. In that case, the running time is $O(n^{O(1)}(\sqrt{3})^n)$.

We need to find a polynomial many initial assignments that cover all possible solutions with distance less than or equal to $n/2$.

Try the following method:

For $j = 1$ to n For $i = j$ to n

f_{ij} = variables j to i equal to 1, and the rest equal to zero $\text{Explore}(\phi, f_{ij}, n/2)$

I think all other possible assignments are within distance $n/2$ of one assignments above. There are n^2 calls to explore function with $d = n/2$. Thus the running time is $O(\text{poly}(n)(\sqrt{3})^{(n/2)})$

7. (Chapter 10 Problem 5 of the text book) 15 points

Given a node weighted graph G (each node v of G is assigned a positive weight $w(v)$), a minimum weight dominating set D of G is a subset of $V(G)$ such that for every node u of G , either $u \in D$ or u is adjacent to a node $v \in D$ and $\sum_{v \in D} w(v)$ is minimized. Give a

polynomial time dynamic programming algorithm (optimal solution structure, Bellman equation, pseudo code, and running time) to find a minimum weight dominating set D and the weight of D in a tree.

Solution:

For each u which is the root of a sub-tree, there are two options.

1. Include it and choose the minimum weight sub-tree rooted at each of its children, no matter include or not includes each child v
2. Not include it and include at least one of its child and the optimal sub-tree of each of the children. The child to include should has the least difference between the optimal solution of sub-tree starting from that node when including v or not including it.

Bellman equations:

let $opt_{in}(u)$ be the weight of min-DS in sub-tree rooted at u , containing u . let C_u be the collection of all children of u .

$$\begin{aligned} opt_{in}(u) &= w(u) + \sum_{v \in C_u} \min\{opt_{in}(v), opt_{out}(v)\} \\ opt_{out}(u) &= opt_{in}(j) + \sum_{v \in C_u \setminus \{j\}} \min\{opt_{in}(v), opt_{out}(v)\} \\ j &= argmin_{v \in C_u} \{opt_{in}(v) - opt_{out}(v)\} \end{aligned}$$

Algorithm:

min-Weight-DS-Tree(T)

Input: A node weighted tree T .

Output: Weight of min-weight DS S in T .

for each node u of T in postorder do

if u is a leaf then $M_{in}[u] = w(u); M_{out}[u] = 0$

for each child v of U :

$j = argmin_{v \in C_u} \{M_{in}(v) - M_{out}(v)\}$

$M_{in}[u] = w(u)$

$S_{in}[u] = \{u\}$

$M_{out}[u] = M_{in}[j]$

$S_{in}[u] = \{j\}$

if $M_{in}(v) < M_{out}(v)$, then

$M_{in}[u] = M_{in}[u] + M_{in}(v)$

$S_{in}[u] = S_{in}[u] \cup \{v\}$

if $v \neq j$ then:

$M_{out}[u] = M_{out}[u] + M_{in}(u)$

$S_{out}[u] = S_{in}[u] \cup \{v\}$

if $M_{out}(v) < M_{in}(v)$, then

$M_{in}[u] = M_{in}[u] + M_{out}(v)$

$S_{in}[u] = S_{out}[u] \cup \{v\}$

if $v \neq j$ then:

```
 $M_{out}[u] = M_{out}[u] + M_{out}(u)$ 
 $S_{out}[u] = S_{out}[u] \cup \{v\}$ 
    end for
end for
if  $M_{in}[r] < M_{out}[r]$ 
    return  $S_{in}[r]$ 
else  $S_{out}[r]$ 
Return  $\max M_{in}[r], M_{out}[r];$ 
```

Running time:

Two loops, one over each node and another over each child. Thus, running time is $O(n + m)$, but since this graph is a tree, the running time is $O(n)$.